

Hauptseminar: Rechnerarchitektur und Programmierung

Evaluation von Visualisierungsmethoden zur skalierbaren
Darstellung von Laufzeit- und Strukturunterschieden in
parallelen Programmabläufen

Ronny Brendel (ronny.brendel@tu-dresden.de)

Tutor: Matthias Weber (matthias.weber@tu-dresden.de)

26th April, 2014

Contents

- Introduction
 - Tracing, Vampir, Challenges
 - Differences, Restrictions
- Ideas & Evaluation
 - Profile
 - Call Matrix
- Conclusion
- Future Work
- Sources

Introduction > Tracing


● Profiling

- Information accumulated per function

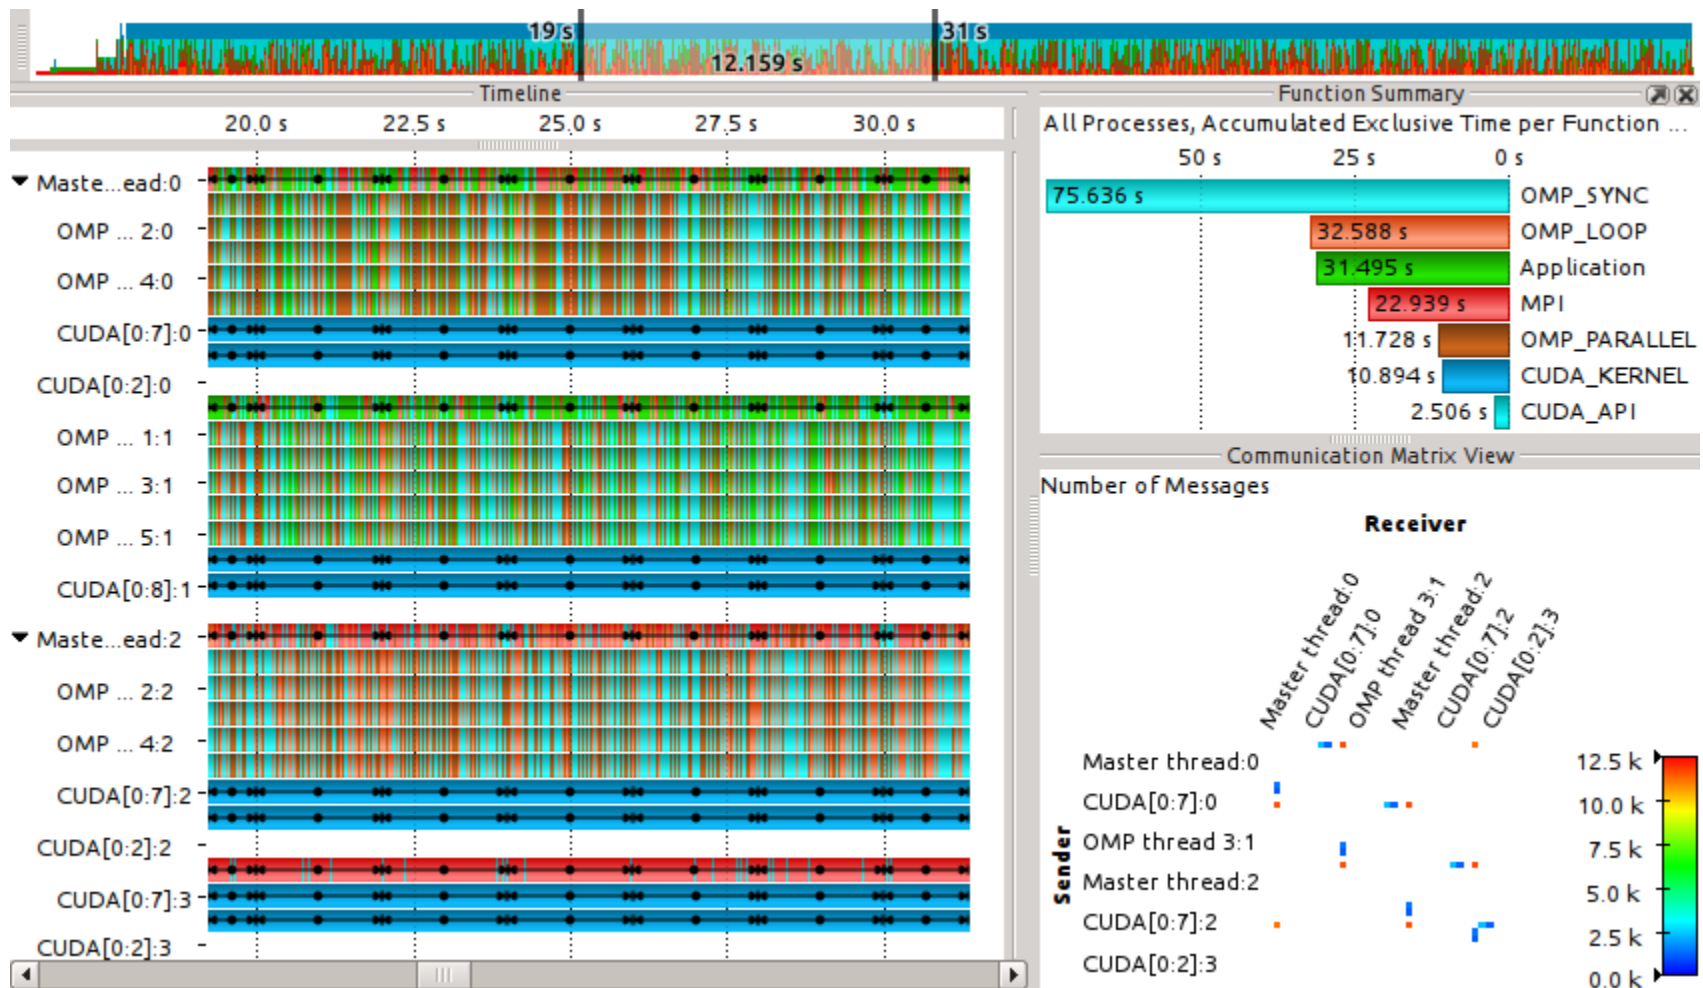
cumulative seconds	self seconds	calls	name
0.31	0.05	25195	QList::isEmpty()
0.40	0.04	30239	QList::Node::t()
0.44	0.04	12294	QList::end()
0.55	0.03	3696	QList::end()
0.70	0.03	4939	handleEnter
0.73	0.03	36939	handleLeave
0.88	0.02	99207	void std::swap()

● Tracing

- Complete information about a program run



Introduction > Vampir



Introduction > Challenges

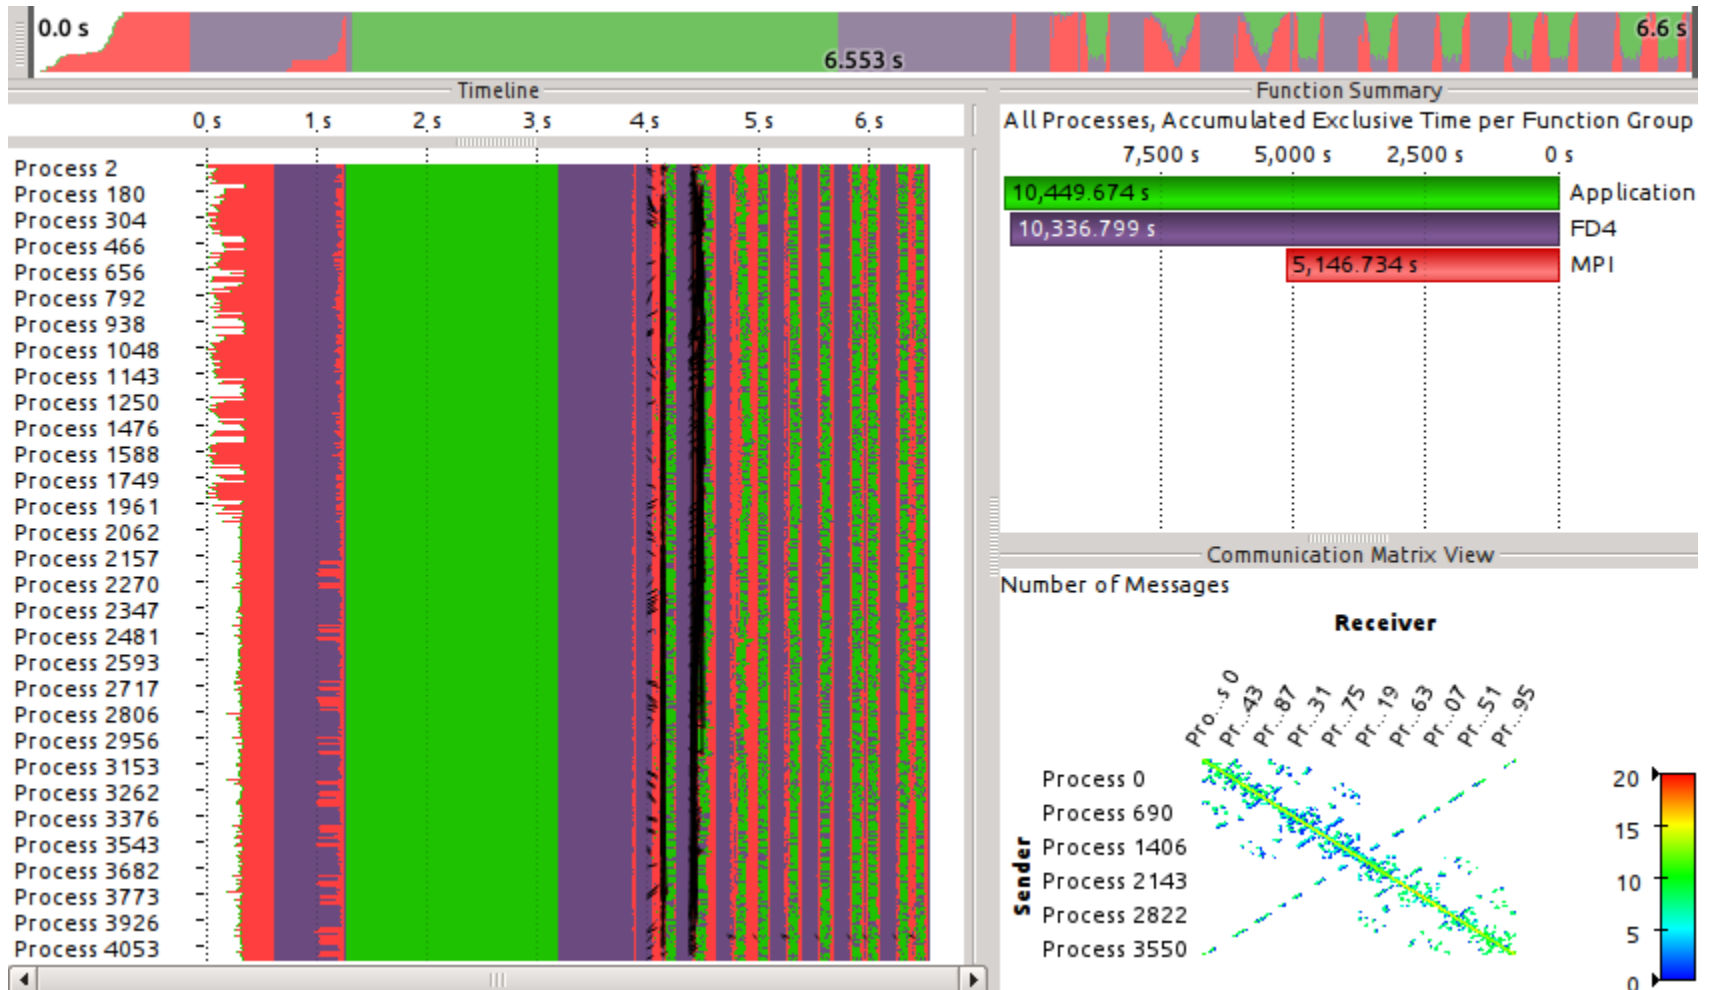
● Datastructures & algorithms:

- Limited main memory size
- Achieving scalability wrt the:
 - Number of processes in the trace
 - Number of processes used for analysis
 - Trace length and detail

● Visualisation:

- Limited number of Pixels on a screen
- Achieving scalability wrt the number of processes in the trace
- Aiding the user to gain insight into his program's behaviour and find performance problems


Introduction > Challenges



Introduction > Differences

Why differences?

- Present new information
 - Visualise timing differences between similar processes
 - Visualise the impact of optimising a program
 - Compare runs of the same program on different platforms
- Improve scalability of existing views
 - Preserve screen real estate by e.g. merging similar processes



- Aid automatic analysis
 - Detect timing differences between structurally similar processes

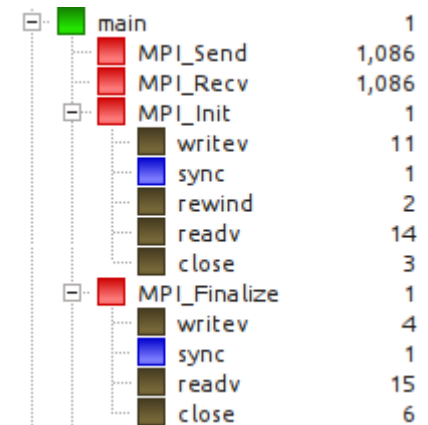

Introduction > Restrictions

- Because visualising differences between traces in general is much ground to cover, we focus on ...
 - ... the function call stack + timing
(no communication, no performance counters, ...)
 - ... comparing processes inside the same program run
 - ... offline analysis
 - ... visualising/comparing profile-ish information

Ideas & Evaluation

Profil-ish information:

- Profile
 - Accumulated information for each function
- Call Tree
 - Accumulated information per call stack configuration
- Call Matrix
 - Accumulated information per caller/callee pair



Ideas & Evaluation

What information are we comparing?



- Execution times

- Exclusive
- Inclusive


- Derived Values

- Min, max, average, standard deviation
- Median, Quartiles/Percentiles
→ Boxplot

- Number of invocations





● Exclusive time, four processes




Ideas & Evaluation > Profile

- Invocation count, four processes
- Exclusive time, four processes



Ideas & Evaluation > Profile



● Exclusive time, 64 processes



Ideas & Evaluation > Profile

Invocation count, 64 processes

Exclusive time, 64 processes





Ideas & Evaluation > Call Matrix

- Color-coded box plot
 - Area is divided in $1 + 3 + 6 + 3 + 1$ parts
 - Linear gradient:




second percentile
first quartile
median
third quartile
98th percentile




Ideas & Evaluation > Call Matrix

- Ping pong example, 1 process, inclusive time




Ideas & Evaluation > Call Matrix

- Ping pong example, 1 process, inclusive time
- main selected





Ideas & Evaluation > Call Matrix

- Ping pong example, 1 process, inclusive time
- main + interesting sub-calls selected




Ideas & Evaluation > Call Matrix

- Ping pong example, 4 processes, inclusive time




Ideas & Evaluation > Call Matrix

Linpack, 32 processes, inclusive time



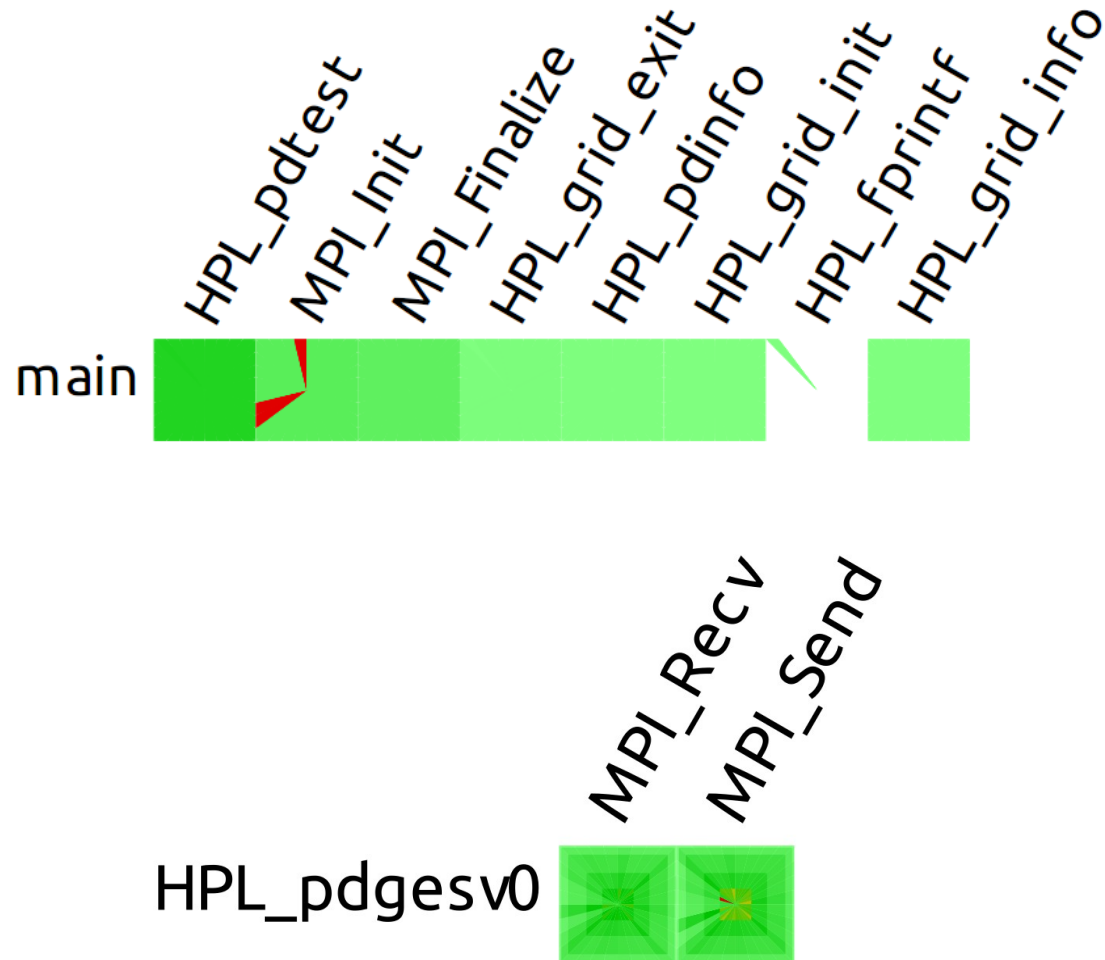
Ideas & Evaluation > Call Matrix

- Linpack, 32 processes, inclusive time



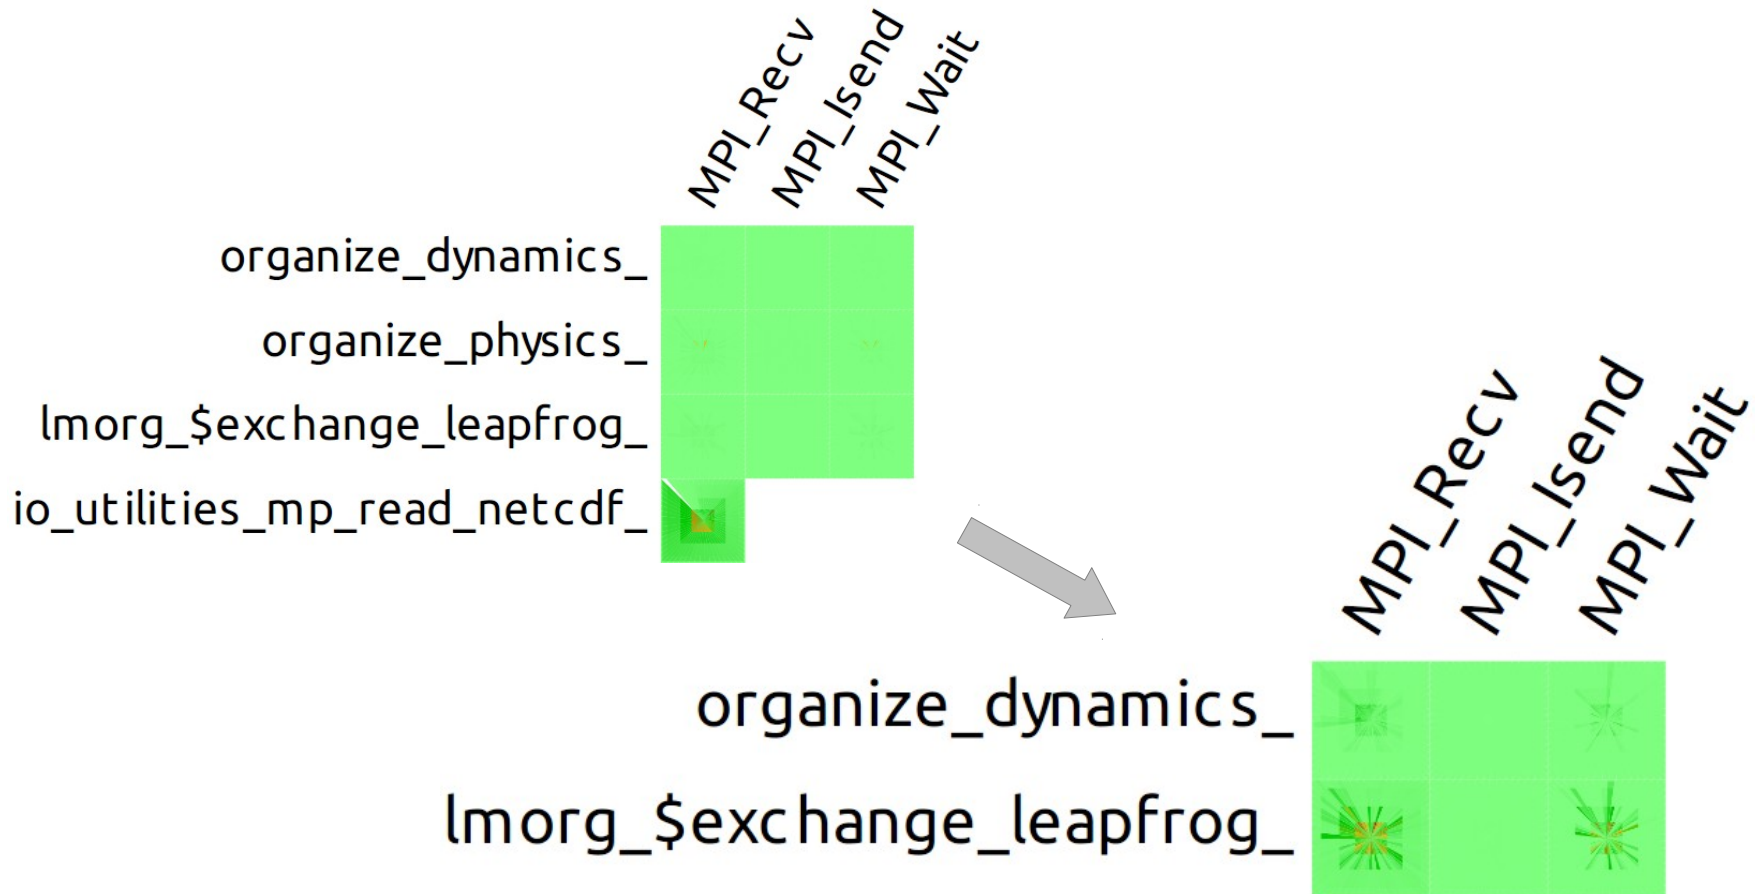
Ideas & Evaluation > Call Matrix

- Linpack, 32 processes, inclusive time



Ideas & Evaluation > Call Matrix

- COSMO-SPECS, 32 processes, exclusive time



Conclusion

- Comparing profiles:
 - Good for getting a qualitative overview
 - Not suitable for finding details
- Comparing call matrixes:
 - Not suitable for visualisation in its current state
 - Improvable through automatic filtering/grouping of functions and processes
 - Simple, scalable and helpful for automatic analysis
- Generally:
 - Boxplots are good
 - Median/Percentiles are good (more difficult to determine, though)
 - Filtering profiles by callers makes sense

Future Work

● Short-term:

- Compare/Visualise call trees
- Develop similarity metrics based on profiles and call matrixes
 - Automatically group processes

● Long-term:

- Visualise differences and similarities in structure and timing of traces (not profiles, but stack over time)

- [1] Andreas Knüpfer, Holger Brunst, Jens Doleschal, Matthias Jurenz, Matthias Lieber, Holger Mickler, Matthias S. Müller, and Wolfgang E. Nagel. The Vampir Performance Analysis Tool-Set. In *Tools for High Performance Computing*, pages 139–155. Springer, 2008.
- [2] Matthias Weber, Ronny Brendel, and Holger Brunst. Trace File Comparison with a Hierarchical Sequence Alignment Algorithm. In *IEEE 10th International Symposium on Parallel and Distributed Processing with Applications*, pages 247–254. IEEE, 2012.
- [3] Matthias Weber, Kathryn Mohror, Martin Schulz, Bronis R. de Supinski, Holger Brunst, and Wolfgang E. Nagel. Alignment-Based Metrics for Trace Comparison. In *Euro-Par 2013 Parallel Processing*, pages 29–40. Springer, 2013.

Thank You!

