



TECHNISCHE
UNIVERSITÄT
DRESDEN

Center for Information Services and High Performance Computing (ZIH)

Interim Diploma Presentation

Automatic Profile-Based Characterization of Performance Traces for Highly Parallel Applications

Ronny Brendel

Professor: Prof. Dr. Wolfgang E. Nagel

Tutors: Matthias Weber, Dr. Holger Brunst

Dresden, 3rd September, 2014



Contents

- Introduction
 - Performance Analysis, Vampir, Challenges
 - Differences in Traces
- Classifying Process Traces
 - Current Approaches
 - Methodology
 - Evaluation
- A Novel Profile Display
 - State of the Art
 - Outlook
- Conclusion, Future Work, References

Introduction > Performance Analysis


● Profiling

- Information accumulated per function and process
- Sampled
- Small overhead

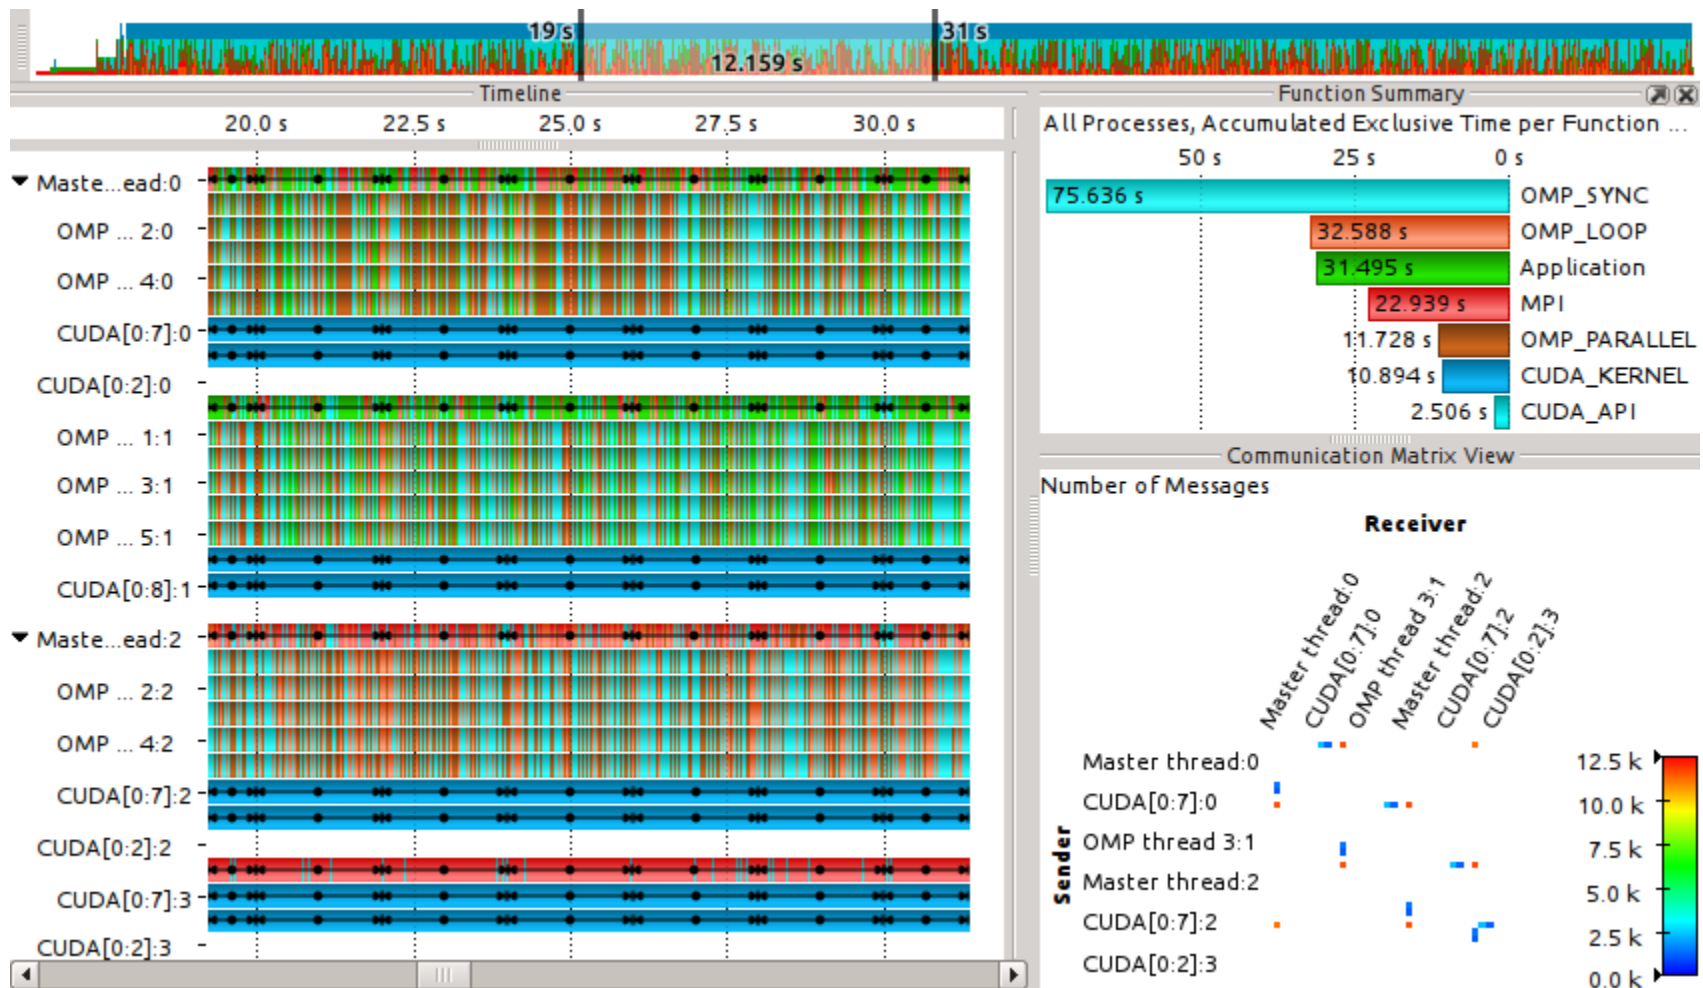
cumulative seconds	self seconds	calls	name
0.31	0.05	25195	QList::isEmpty()
0.40	0.04	30239	QList::Node::t()
0.44	0.04	12294	QList::end()
0.55	0.03	3696	QList::end()
0.70	0.03	4939	handleEnter
0.73	0.03	36939	handleLeave
0.88	0.02	99207	void std::swap()

● Tracing

- Complete information about a program run
- Instrumented
- Large overhead



Introduction > Vampir



Introduction > Challenges

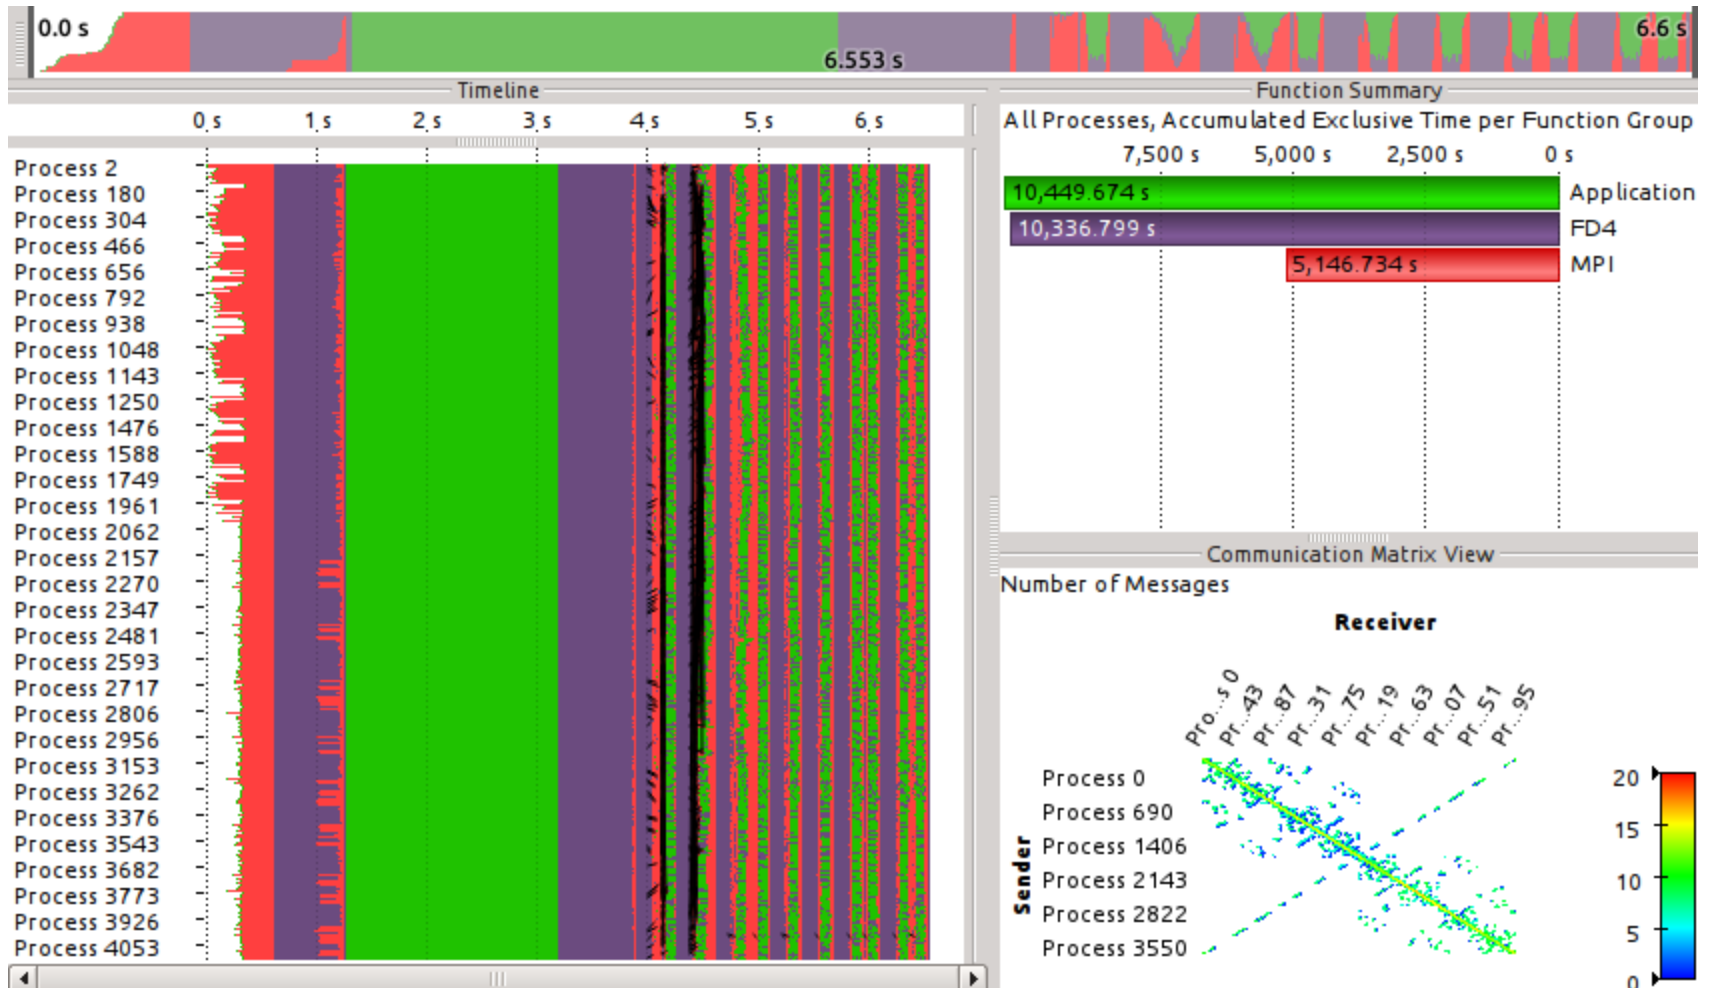
■ Datastructures & algorithms:

- Limited main memory size
- Achieving scalability wrt:
 - Number of processes in the trace
 - Number of processes used for analysis
 - Tracefile Size

■ Visualisation:


- Limited number of Pixels on a screen
- Achieving scalability wrt the number of processes in the trace

Introduction > Challenges



Introduction > Differences in Traces

- How does detecting differences between process traces aid performance analysis?
 - Present new information
 - Visualise timing differences between similar processes
 - Visualise the impact of optimisation
 - Compare runs of the same program on different platforms
 - Improve scalability of existing views
 - Preserve screen real estate by e.g. merging similar processes
 - Aid automatic analysis
 - Detect timing differences between structurally similar processes




Classifying Process Traces > Current Approaches

- Take any similarity measure. Calculating it for all pairs of processes already requires an all-to-all comparison, which is not feasible for 50k+ processes.
- Clustering n processes using common clustering algorithms, e.g. k-means and DBSCAN, takes at least n^2 steps
- Some algorithms require a predefined number of clusters or make assumptions about the data layout, which leads to unnatural clusters.
- To handle large amounts of processes properly, we need something better

Classifying Process Traces > Methodology

- A simple, structural similarity metric
 - Each process p is assigned a set of functions $\{A \mid A \text{ is called on } p \text{ at least once}\}$
 - $\text{Similarity}(P1, P2) := |P1 \cap P2| / |P1 \cup P2|$
 - Alternatively: each process p is assigned a set of function pairs $\{(A, B) : A \text{ calls } B \text{ on process } p \text{ at least once}\}$
 - Similarity is calculated analogously






- Next step: Calculate and store similarity between processes

Classifying Process Traces > Methodology

Use a concept lattice


- Total order, e.g. (\mathbb{N}, \leq)
- Partial order, e.g. $(\mathbb{N}, |)$
- Bounded lattice, e.g. $(\{A, B\}, \subseteq)$



- Concept lattice $(\mathcal{P}, \mathcal{F}, \mathcal{I} \subseteq \mathcal{P} \times \mathcal{F})$

where e.g. $\mathcal{P} := \{P1, P2, P3\}$, $\mathcal{F} := \{A, B, C\}$,

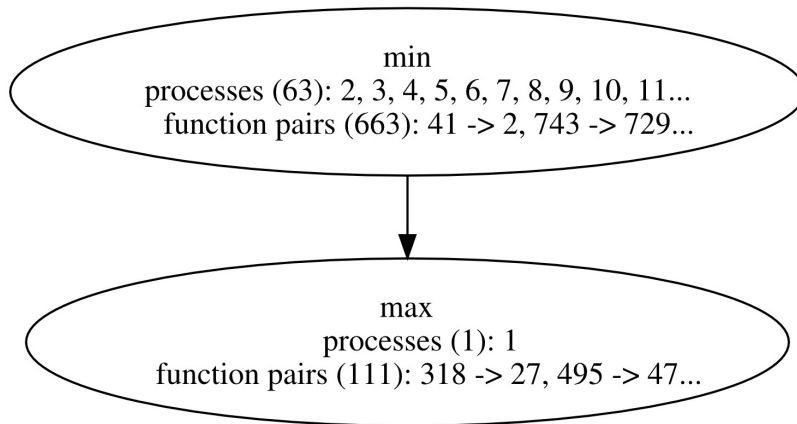
$\mathcal{I} := \{ \quad P1 : A \quad P2 : B, C \quad P3 : A, B, C \quad P4 : B, C \quad \}$



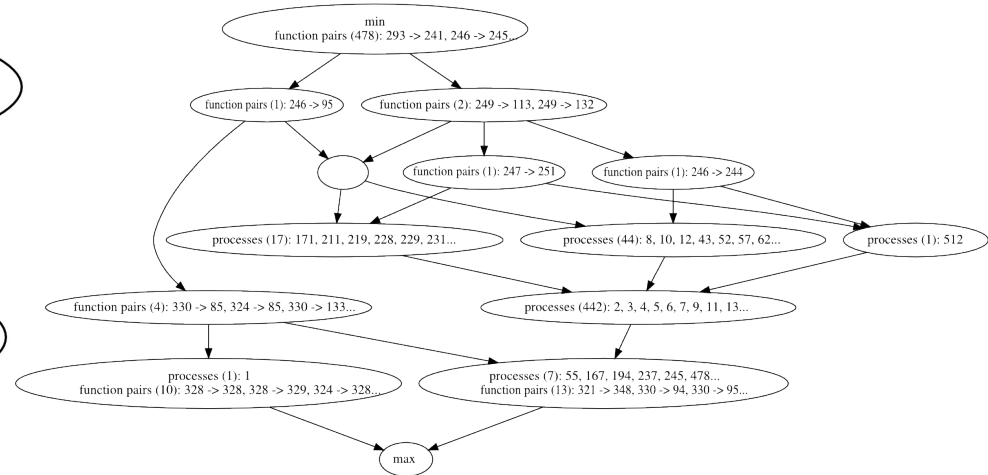
Classifying Process Traces > Methodology

- The result is a directed acyclic graph of „function pair inheritance“
- Expected complexity for building and storing the graph is linear, worst case exponential

WRF 64 processes



AMG 512 processes

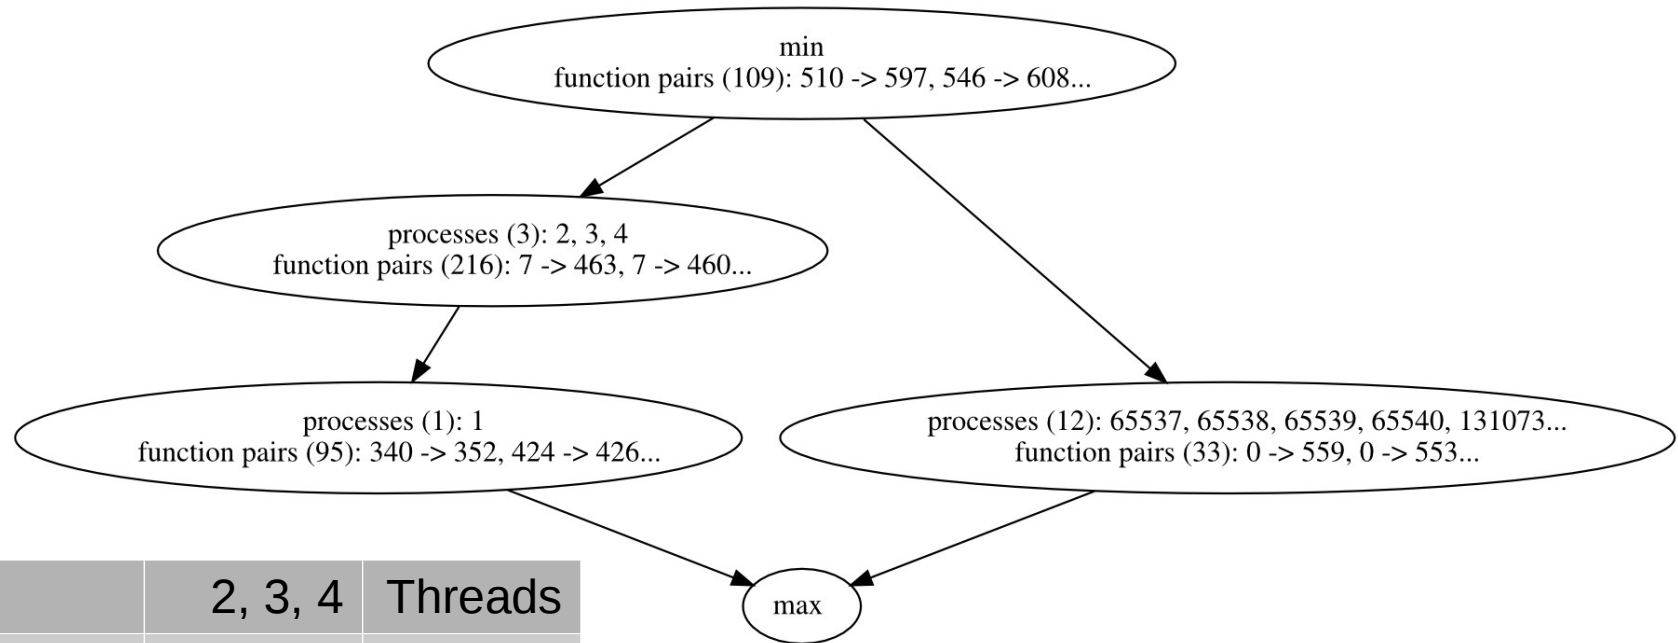


- Next step: calculate similarity between „equivalence classes“

Classifying Process Traces > Methodology

■ $\text{Similarity}(P1, P2) := |P1 \cap P2| / |P1 \cup P2|$


- Simple addition and division of set sizes. No set operations involved.
- Similarity matrix can then be calculated in $O(g^2n)$ steps



	2, 3, 4	Threads
1	77%	24%
2, 3, 4		30%

Classifying Process Traces > Methodology


- Slightly modifying this technique lets us find out whether a process calls the same and more functions than another. This way we can:
 - Detect inlining
 - Detect whether processes are dissimilar because one performs additional calls, e.g. logs output
 - Detect different levels of coarseness of trace recording



$P1 = \{(main, A), (A, B)\}$


$P2 = \{(main, B)\}$

Similarity = 0%




Classifying Process Traces > Methodology

- 1st step: Calculate transitive closure of both P and Q



$P1 = \{(main, A), (A, B)\} \cup \{(main, B)\}$
 $P2 = \{(main, B)\}$



- 2nd step: Build the concept lattice
- 3rd step: $\text{SubsumptionMetric}(P1, P2) := |P1 \cap P2| / |P2|$
- Example result is 100%





	P1	P2
P1		1
P2	1/3	

Classifying Process Traces > Evaluation

Trace			Result		
Application	Proc.	Size (MiB)	Construction Time (%)	Nodes	Groups
LINPACK	832	145	≤ 0.1	3	3
Gromacs	36	3,700	≤ 0.1	24	11
COSMO-SPECS	100	976	≤ 0.1	1	1
FD4	4096	166	≤ 0.1	2	2
WRF	64	283	≤ 0.1	2	2
HOMME	1024	179	≤ 0.1	3	3
LULESH	432	354	0.27	182	35
PICongGPU	39	286	≤ 0.1	60	17
BT	16	150	≤ 0.1	5	3

Classifying Process Traces > Evaluation


AMG2006




Classifying Process Traces > Evaluation

● ParaDiS

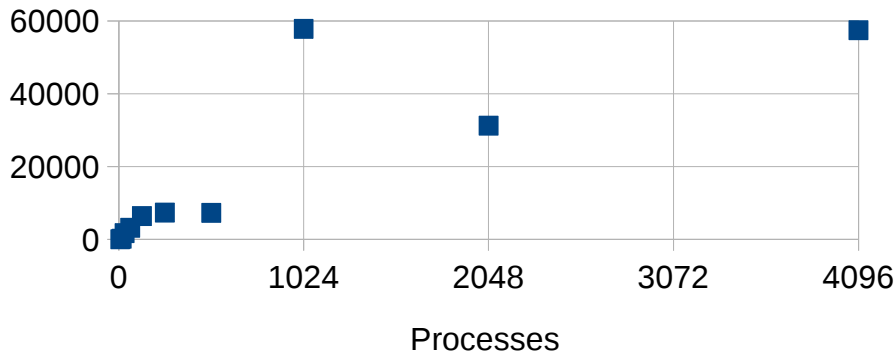
Construction Time




Tracefile Loading Time



Number of Nodes in the Lattice

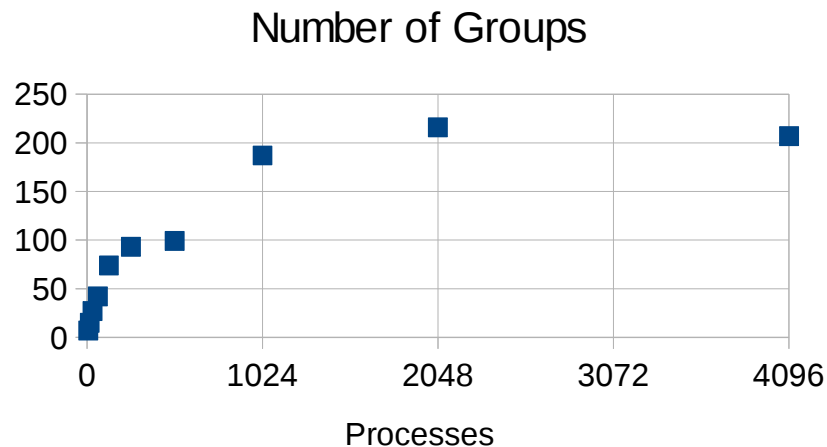




Construction Time div. Loading Time



Classifying Process Traces > Evaluation


ParaDiS



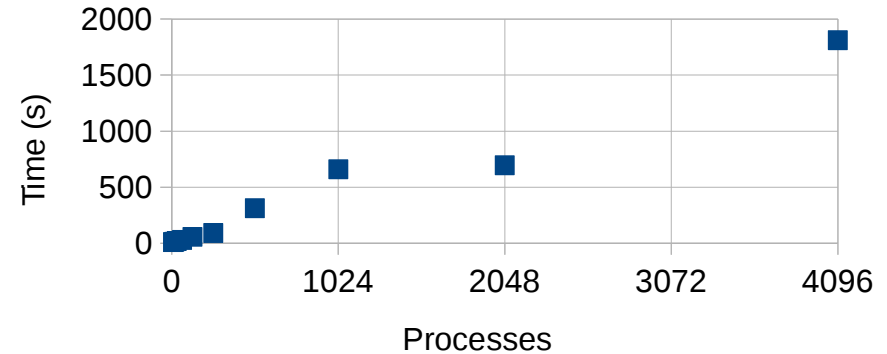
Classifying Process Traces > Evaluation

- ParaDiS, using called functions instead of function pairs

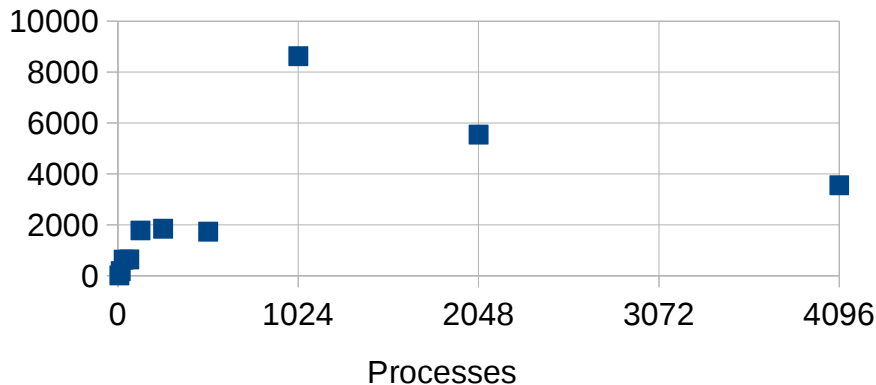
Construction Time




Tracefile Loading Time



Number of Nodes in the Lattice

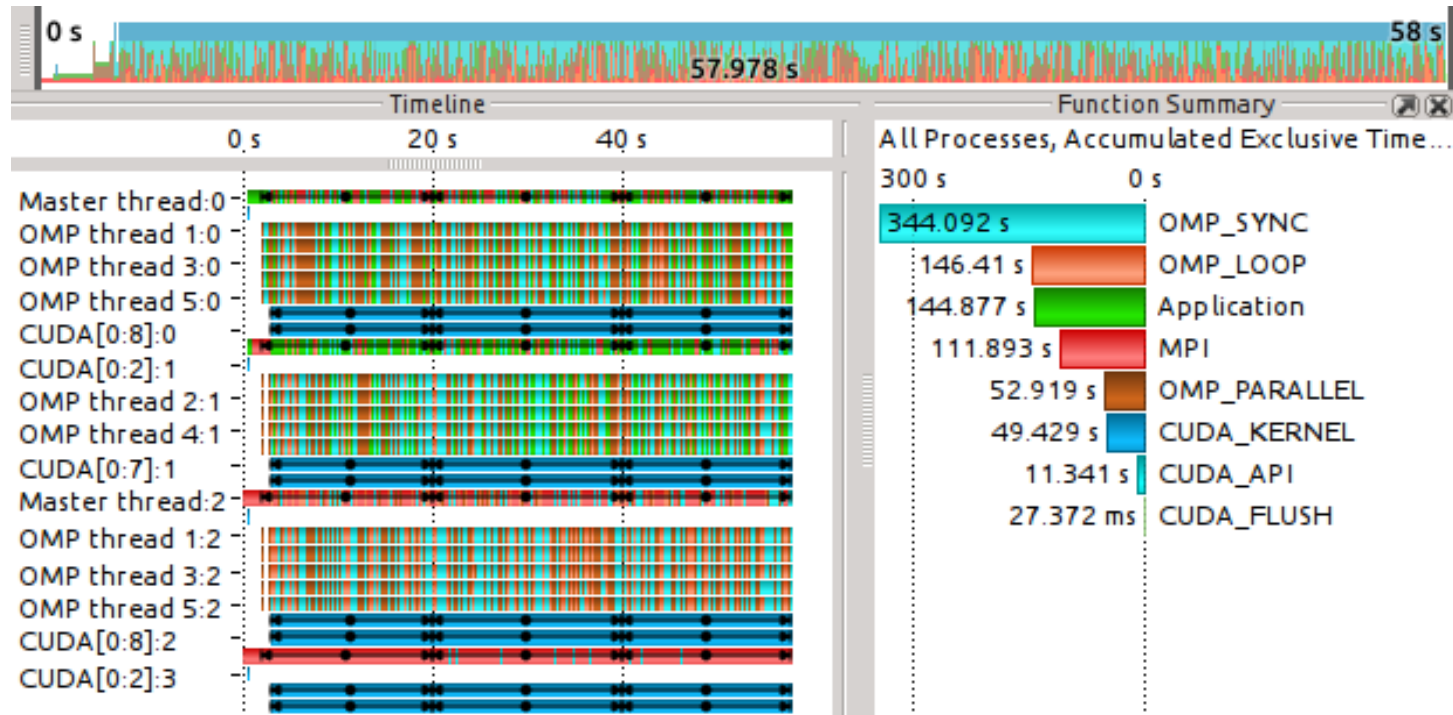


Construction Time div. Loading Time



Classifying Process Traces > Evaluation

- Gromacs
 - Decreasing tracing detail from Masterthread 0 to 3



Classifying Process Traces > Evaluation

● Gromacs

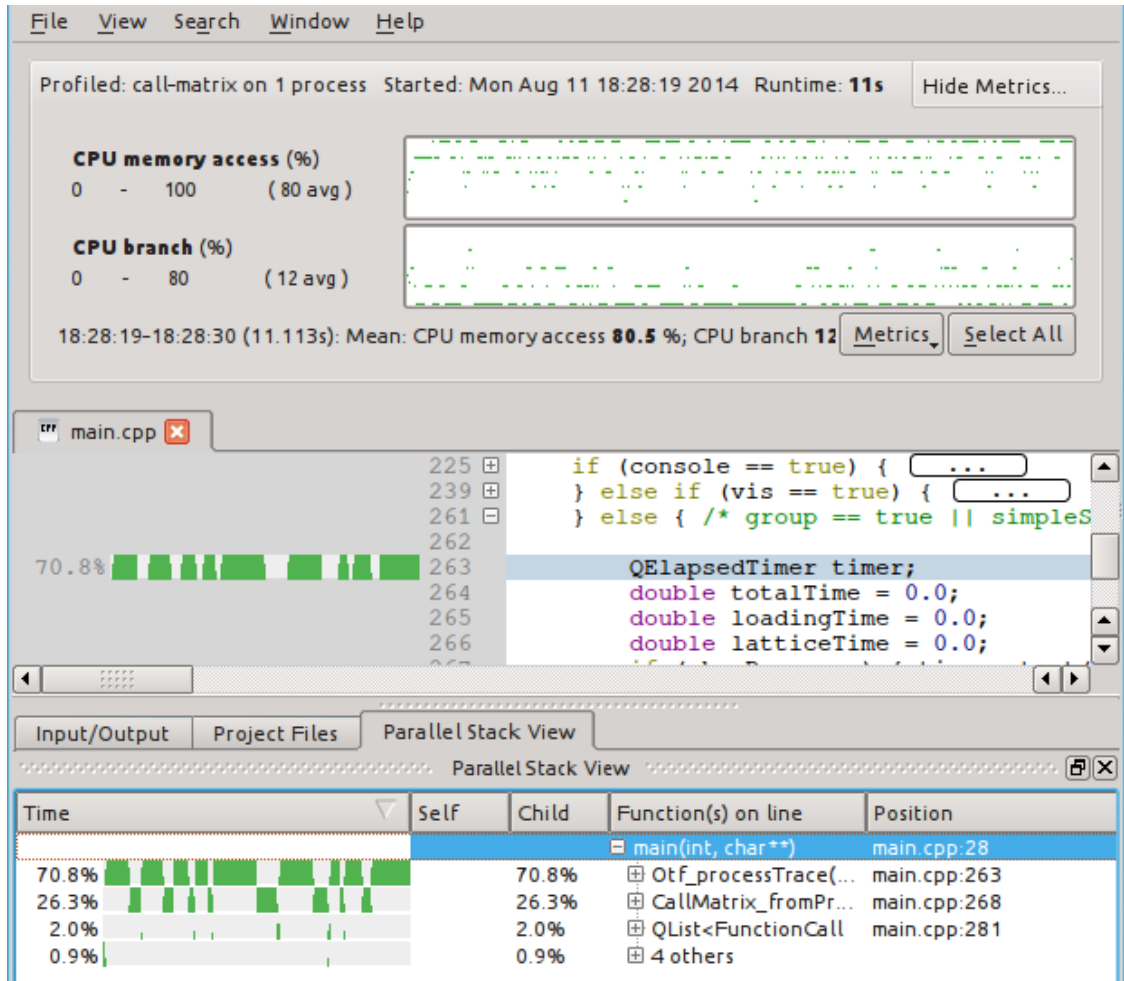
- Subsumption metric matrix, function appearance (not call relations), for all master threads

	0	1	2	3
0		100%	99%	98%
1	76%		99%	98%
2	24%	31%		100%
3	8%	10%	33%	

A Novel Profile Display > State of the Art

● Allinea Map

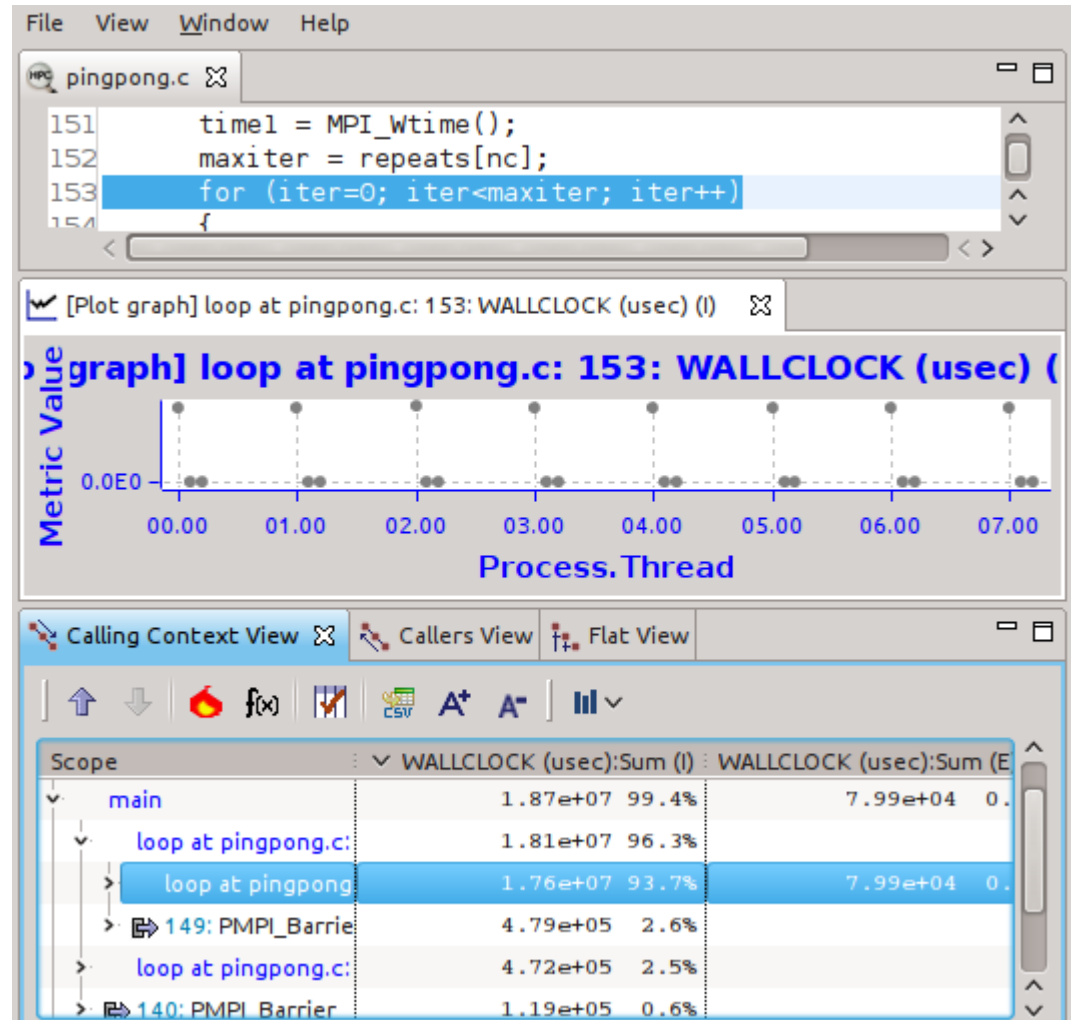
- Call path profiling
- Accumulates over all processes
- Sliced information (vampir-esque)
- Various metrics
- Very intuitive and simple GUI



A Novel Profile Display > State of the Art

HPCToolkit

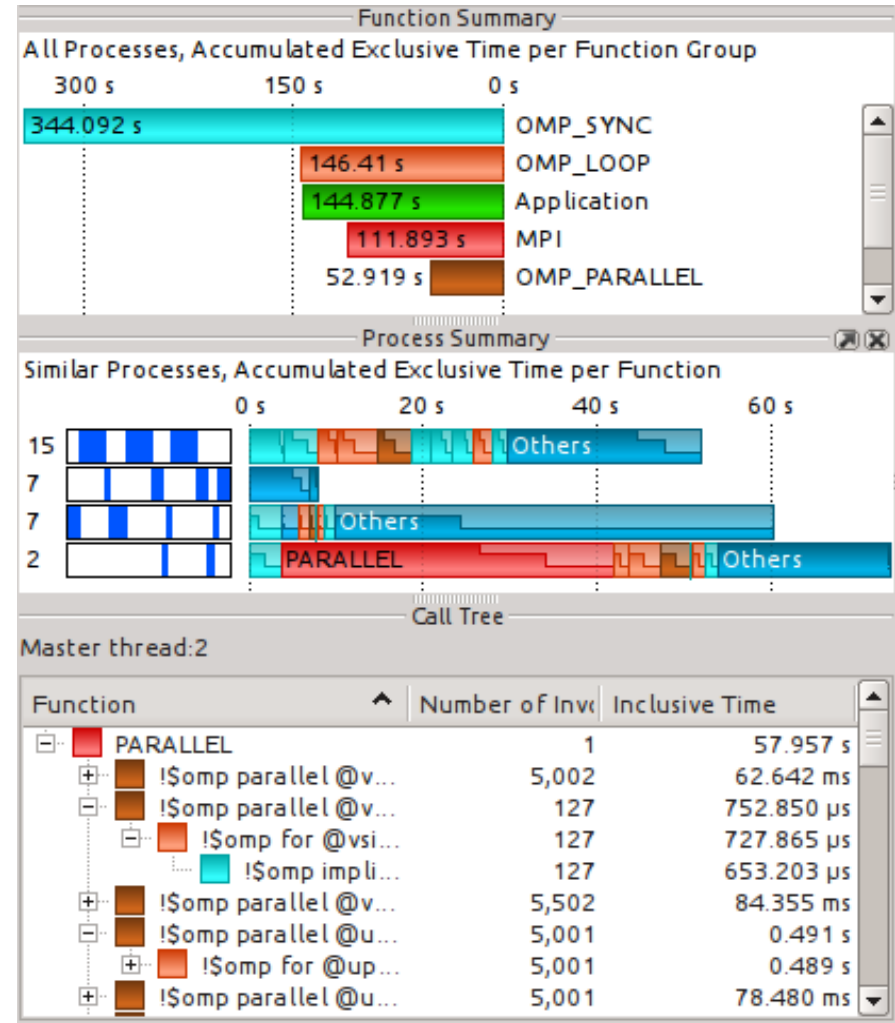
- Call path profiling
- Accumulates over all processes
- Plots for one function value x processes
- Custom metrics
- Usability is ok



A Novel Profile Display > State of the Art

● Vampir

- Powerful profile displays
- Supports clustering similar profiles
- Usability is good, but needs practice/knowledge



A Novel Profile Display > Outlook

- What can we improve?
 - Replace the Function Summary, Call Tree and Process Summary displays in Vampir with one unified display.
 - Usability
 - Comparative analysis

A Novel Profile Display > Outlook

● Ideas: General Usability

- Fewer bars, maybe none
- Improved search functionality
- Proper undo stack
- Context menu that is actually context-sensitive
- ...

A Novel Profile Display > Outlook

● Ideas:

- Show useful and simple information at first sight
- Enable easy regrouping of functions
- Highlighting points of interest
 - Special sorting options
 - ?
- Display timing variations between different calls or processes using e.g. box plots
- Comparative Analysis:
 - Filter processes using the new clustering
 - ?

Conclusion

- Introduced a structural similarity metric
- Developed methods to cluster and compare processes in $< O(n^2)$ steps
 - works with any set-based similarity metric
- First steps towards a better profile viewer

Future Work

- Further scalability testing of the clustering
- Develop a novel, scalable profile viewer
- Use the developed methods in an actual visualisation

References

- [1] Andreas Knüpfer, Holger Brunst, Jens Doleschal, Matthias Jurenz, Matthias Lieber, Holger Mickler, Matthias S. Müller, and Wolfgang E. Nagel. The Vampir Performance Analysis Tool-Set. In *Tools for High Performance Computing*, pages 139–155. Springer, 2008.
- [2] Matthias Weber, Kathryn Mohror, Martin Schulz, Bronis R. de Supinski, Holger Brunst, and Wolfgang E. Nagel. Alignment-Based Metrics for Trace Comparison. In *Euro-Par 2013 Parallel Processing*, pages 29–40. Springer, 2013.
- [3] Dean Van Der Merwe, Sergei Obiedkov, and Derrick Kourie. Addintent: A new incremental algorithm for constructing concept lattices. In *Concept Lattices*, pages 372–385. Springer, 2004.

References

- [4] Robert D Falgout and Ulrike Meier Yang. hypre: A library of high performance preconditioners. In Computational Science—ICCS 2002, pages 632–641. Springer, 2002.
- [5] ParaDiS. <http://paradis.stanford.edu>, 2014-08-27.
- [6] Allinea Map. <http://www.allinea.com/products/map>, 2014-08-27.
- [7] HPCToolkit. <http://hpctoolkit.org>, 2014-08-27.

Thank You!